

Task 1 Group 6 README

Accessing the dataset

Logging into cluster

If you are off campus, a VPN must be used. The client is often Cisco. To log into the cluster, the following command is used:

```
ssh YOURusername@hadoop-nn001.cs.okstate.edu
```

If you are on Linux, you can use `sshpass` to skip typing in your password each time. It may be possible to use this on Windows, though the author of this does not use Windows and is unaware if its possible. The discussion of this method of logging in is located at the stack overflow link

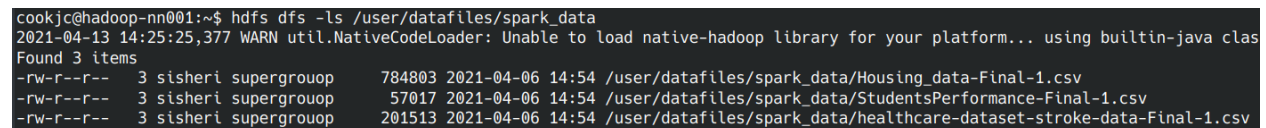
```
# https://stackoverflow.com/a/16734873/11637415
```

```
sshpass -p "YOURPassword" ssh -o \
StrictHostKeyChecking=no YOURusername@hadoop-nn001.cs.okstate.edu
```

Moving dataset to local machine

The dataset is located in the hadoop file system (HDFS) on the cluster; `/user/datafiles/spark_data`. To list the files in that path the following command can be typed:

```
hdfs dfs -ls /user/datafiles/spark_data
```



```
cookjc@hadoop-nn001:~$ hdfs dfs -ls /user/datafiles/spark_data
2021-04-13 14:25:25,377 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java clas
Found 3 items
-rw-r--r--  3 sisheri supergroup    784803 2021-04-06 14:54 /user/datafiles/spark_data/Housing_data-Final-1.csv
-rw-r--r--  3 sisheri supergroup    57017 2021-04-06 14:54 /user/datafiles/spark_data/StudentsPerformance-Final-1.csv
-rw-r--r--  3 sisheri supergroup   201513 2021-04-06 14:54 /user/datafiles/spark_data/healthcare-dataset-stroke-data-Final-1.csv
```

Figure 1: The hdfs files on the cluster for the big data group assignment

The `Housing` dataset has been assigned to us (Group 6). Spark has a `write once, run everywhere` methodology. We now want to move the `Housing_data-Final-1.csv` dataset to our `/home` directory so we can then copy it off the cluster to our own local machine for running spark programs.

```
hadoop fs -copyToLocal \
/user/datafiles/spark_data/Housing_data-Final-1.csv .
```

Now the file can be moved off the cluster to our local machine. From a terminal window that is not logged into the cluster, `scp` can now be used to move the data off the cluster.

```
scp -r \
cookjc@hadoop-nn001.cs.okstate.edu:/home/cookjc/\
Housing_data-Final-1.csv .
```

Description of dataset

The spark program is used to determine the number of features (columns) and rows that are in the dataset. The schema of the dataset is also displayed. The schema is an overview of the column names and the variable types that are in each column (or feature).

Figure 2: The number of rows and columns and the schema of the dataframe

At the very bottom of the data set (row 11,583) there exists a description of what some of the keys in the dataset mean:

- Suburb - name of suburb
- Type:
 - h - house, cottage, villa, semi, terrace

```

The number of rows in the dataset: 11597
The number of columns in the dataframe: 11
root
|-- Suburb: string (nullable = true)
|-- Type: string (nullable = true)
|-- Price: integer (nullable = true)
|-- Distance: double (nullable = true)
|-- Zipcode: integer (nullable = true)
|-- #Bedroom: integer (nullable = true)
|-- #Bathroom: integer (nullable = true)
|-- #-Car Garage: integer (nullable = true)
|-- Lot_size: integer (nullable = true)
|-- Region_name: string (nullable = true)
|-- Property_count: integer (nullable = true)

```

Figure 1: F2_DatasetDescription

- u - unit, duplex
- t - townhouse site, development side
- o - res, other residential
- Price: Price in dollars (\$)
- Distance: Distance from CBD
- Zipcode: the zipcode where the unit is located
- Bedroom: number of bedrooms
- Bathroom: number of bathrooms
- Car Garage: number of car garages
- Lot Size: the size of the property (in acres)
- Region name: General region (West, North West, North, North East, etc.)
- Property count: Number of properties that exist in the suburb

There are 7 region names, 3 housing types and 312 suburbs.

Issues with the data

There are rows at the bottom at the data that helped to describe the data, but it did not contain any meaningful data for processing and has been removed.

Here is a list of the csv files and a description of what they are:

- Housing_data-Final-1.csv - the original csv supplied
- Housing_data-Final-2.csv - the bottom 20 rows of garbage data are removed